

UrbanCart — Ecommerce Analytics

Executive Summary

UrbanCart's operational data (2,500 customers, 9,000 orders across 34 months) was cleaned, reconciled with two legacy data sources, and analyzed to answer four questions: who are our best customers, what will revenue look like next year, where is our data unreliable, and where does the evidence run out.

What we found

- Only 57% of orders are completed — the rest are cancelled, returned, or still pending. This is worth an operational look independent of this analysis.
- We found and fixed real data problems: 48 duplicate customer records, 186 duplicate order lines, and — the biggest one — 253 of 255 product SKUs (99%) where our internal price doesn't match the supplier catalog, by an average of \$9.44. Recommend a merchandising review before either price list is trusted.
- Customers split into 6 value tiers. The average customer is worth about \$1,916.85 over their lifetime, but our top segment ("Champions") is worth \$3,257.53 on average — nearly 2x the mean.
- A group of past high-spenders has gone quiet ("At Risk (High Value)") — this is the single best reactivation target we identified.
- Next 12 months of revenue is likely to fall between \$2.13M and \$2.98M, with \$2.55M as the most likely outcome, based on 10,000 simulations of historical revenue patterns.
- We looked for a link between price and how much customers buy, and between products purchased together — neither showed a real pattern in this data. We're not recommending any pricing or bundling changes based on these two analyses; the honest answer is "not enough evidence yet."

What we recommend, in order

#	Action
1	Launch a win-back campaign for the "At Risk (High Value)" segment (210 customers, ~\$2,734.09 avg. value each).
2	Reconcile the catalog vs. internal pricing discrepancy with merchandising before it affects a pricing decision.
3	Investigate why order completion is only ~57% — a checkout, fulfillment, or returns-policy review.
4	Plan next year's revenue around the \$2.13M–\$2.98M range, not a single number.

A note on the source data: a separate 'Master Specification' document supplied with this data made infrastructure and performance claims that don't hold up against what's actually in the database. This summary and the full report behind it are built only from what the data itself supports.